

# Supervised Learning I

Christoph Kern

Mannheim Machine Learning Modules

*c.kern@uni-mannheim.de*

October 20, 2017



# Outline

- 1 Introduction
- 2 ML Basics
  - Training and test error
  - Validation set, test set, CV
  - Performance measures
- 3 CART
  - Tree growing
  - Tree pruning
- 4 References

# Introduction

## Machine Learning

- Algorithms based on statistical criteria which focus on making predictions based on a data-driven learning process
- “Machine Learning is the field of scientific study that concentrates on induction algorithms and on other algorithms that can be said to “learn”.” (Kohavi / Provost 1998)
- Combines Computer Science and Statistics

## Statistical Learning

- Machine Learning from a “statistical perspective”

# ML Basics

## Unsupervised Learning

- Finding patterns in data using a set of input variables  $X$

## Supervised Learning

- Predicting an output variable  $Y$  based on a set of input variables  $X$ 
  - 1 Learn the relationship between input and output using **training data** (with  $X$  and  $Y$ )

$$Y = f(X) + \varepsilon$$

- 2 Predict the output based on the prediction model (of step 1) for **new test data** (~only  $X$  available)
- continuous  $Y$ : regression, categorical  $Y$ : classification
  - Focus on **prediction** ( $\neq$  causation)

# Training and test error

## Training error

$$\overline{\text{err}} = \frac{1}{n} \sum_{i=1}^n L(y_i, \hat{f}(x_i))$$

- Prediction error based on **training data**

## Test error

$$\text{Err}_{\mathcal{T}} = \mathbb{E}(L(Y, \hat{f}(X)) | \mathcal{T})$$


- Prediction error using **test data** (given training data  $\mathcal{T}$ )

## Expected test error decomposition

$$\text{Err}(x_0) = \text{Bias}^2(\hat{f}(x_0)) + \text{Var}(\hat{f}(x_0)) + \text{Var}(\varepsilon)$$

- Minimizing the expected test error
  - Low bias (deviation between  $\mathbb{E}(\hat{f}(x_0))$  and  $f(x_0)$ ) **and**
  - Low variance ( $\text{Var}(\hat{f}(x_0))$  using different training data)

Figure: Bias-Variance Trade-Off: Training error and test error by model complexity



Hastie et al. 2009

# Validation set, test set, CV

## Validation set approach

- Training set & validation set

- 1 Fit model using one part of training data
- 2 Compute test error for the excluded section

### → Model assessment

- Training set, validation set & test set

- 1 Fit models using training part of training data
- 2 Choose best model using validation set
- 3 Evaluate final model using test set

### → Model tuning & assessment

## Cross-Validation

- LOOCV (Leave-One-Out Cross-Validation)

- 1 Fit model on training data while excluding one case
- 2 Compute test error for the excluded case
- 3 Repeat step 1 & 2  $n$  times

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^n L(y_i, \hat{f}^{-\kappa(i)}(x_i))$$

- k-Fold Cross-Validation

- 1 Fit model on training data while excluding one group
- 2 Compute test error for the excluded group
- 3 Repeat step 1 & 2  $k$  times (e.g.  $k = 5$ ,  $k = 10$ )

$$CV_{(k)} = \frac{1}{k} \sum_{i=1}^k L(y_i, \hat{f}^{-\kappa(i)}(x_i))$$

- Outlook: nested CV, repeated CV, ...



# Performance measures

Performance metrics for regression problems

$$r^2 = 1 - RSS/TSS$$

Root mean squared error (RMSE):

$$\sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2}$$

Median of squared errors (MEDSE):

$$\text{median}((y_1 - \hat{f}(x_1))^2, \dots, (y_n - \hat{f}(x_n))^2)$$

Mean of absolute errors (MAE):

$$\frac{1}{n} \sum_{i=1}^n |(y_i - \hat{f}(x_i))|$$

# Performance metrics for classification

Table: Confusion matrix

		Prediction		
		0	1	
Reference	0	True Negatives (TN)	False Positives (FP)	N'
	1	False Negatives (FN)	True Positives (TP)	P'
		N	P	

Accuracy:  $\frac{TP+TN}{TP+FP+TN+FN}$

Sensitivity / Recall:  $\frac{TP}{TP+FN}$

Specificity:  $\frac{TN}{TN+FP}$

Positive predictive value / Precision:  $\frac{TP}{TP+FP}$

Negative predictive value:  $\frac{TN}{TN+FN}$


Outlook:  $F_1$ ,  $\kappa$ , AUC-ROC, AUC-PR, ...

# CART

## Classification and Regression Trees (CART)

- Approach for partitioning the predictor space into smaller subregions via "recursive binary splitting"
- Results in a "top-down" tree structure with...
  - Internal nodes within the tree
  - Terminal nodes as endpoints
- Can be applied to regression and classification problems
- Can be used as base learner for ensemble methods

Figure: A small tree



James et al. 2013

# Tree growing

Growing a regression tree

Define pairs of regions for all  $X_1, X_2, \dots, X_p$  predictors and cutpoints  $c$

$$\tau_L(j, c) = \{X | X_j < c\} \text{ and } \tau_R(j, c) = \{X | X_j \geq c\}$$

Find split  $s$  which maximizes the reduction in  $RSS$

$$\Delta RSS(s, \tau) = RSS(\tau) - RSS(\tau_L) - RSS(\tau_R)$$

$$RSS(\tau) = \sum_{i \in \tau} (y_i - \hat{y})^2$$

with  $\hat{y}$  being the mean of  $y$  in node  $\tau_m$

## Growing a classification tree

Define pairs of regions for all  $X_1, X_2, \dots, X_p$  predictors and cutpoints  $c$

$$\tau_L(j, c) = \{X | X_j < c\} \text{ and } \tau_R(j, c) = \{X | X_j \geq c\}$$

Find split  $s$  which maximizes the reduction in node impurity

$$\Delta I(s, \tau) = I(\tau) - p(\tau_L)I(\tau_L) - p(\tau_R)I(\tau_R)$$

Impurity measures

$$I_{Gini}(\tau) = \sum_{k=1}^K \hat{p}_k(1 - \hat{p}_k)$$

$$I_{entropy}(\tau) = - \sum_{k=1}^K \hat{p}_k \log \hat{p}_k$$

with  $\hat{p}_k$  being the proportion of observations from class  $k$  in node  $\tau$

---

**Algorithm 1:** Tree growing process

---

```
1 Define stopping criteria;
2 Assign training data to root node;
3 if stopping criterion is reached then
4   |   end splitting;
5 else
6   |   find the optimal split point;
7   |   split node into two subnodes at this split point;
8   |   for each node of the current tree do
9     |   continue tree growing process;
10  |   end
11 end
```

---

A given tree

$$\mathcal{T} = \sum_{m=1}^M \gamma_m \cdot \mathbf{1}_{(i \in \tau_m)}$$

consists of a set of  $m = 1, 2, \dots, M$  nodes which can be used for prediction by...

- Regression
  - ...using the mean of  $y$  for training observations in  $\tau_m$
- Classification
  - ...going with the majority class in  $\tau_m$

→ Prediction surface: Block-wise relationship between features and outcome



# Tree pruning

## Stopping rules

- Minimum number of cases in terminal nodes
- Decrease in impurity exceeds some threshold

→ However, worthless splits can be followed by good splits

## Cost complexity pruning

$$R_{\alpha}(\mathcal{T}) = R(\mathcal{T}) + \alpha|\mathcal{T}|$$

- Find the best subtree by balancing tree quality  $R(\mathcal{T})$  and complexity  $|\mathcal{T}|$
- $\alpha$  controls the penalty on the number of terminal nodes
- $\alpha$  can be chosen through CV

# References

- James, G., Witten, D., Hastie, T., Tibshirani, R. (2013). *An Introduction to Statistical Learning*. New York, NY: Springer.
- Kohavi, R., Provost, F. (1998). Glossary of Terms. *Machine Learning*, 30(2): 271–274.
- Hastie, T., Tibshirani, R., Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. New York, NY: Springer.